

Create tables in Hive and write queries to access the data in the table

AIM:

To create tables in Hive and write queries to access the data in the table.

PROCEDURE:

Installation of Hive:

1. Installing Apache Derby

Install Apache Derby 10.14.2.0

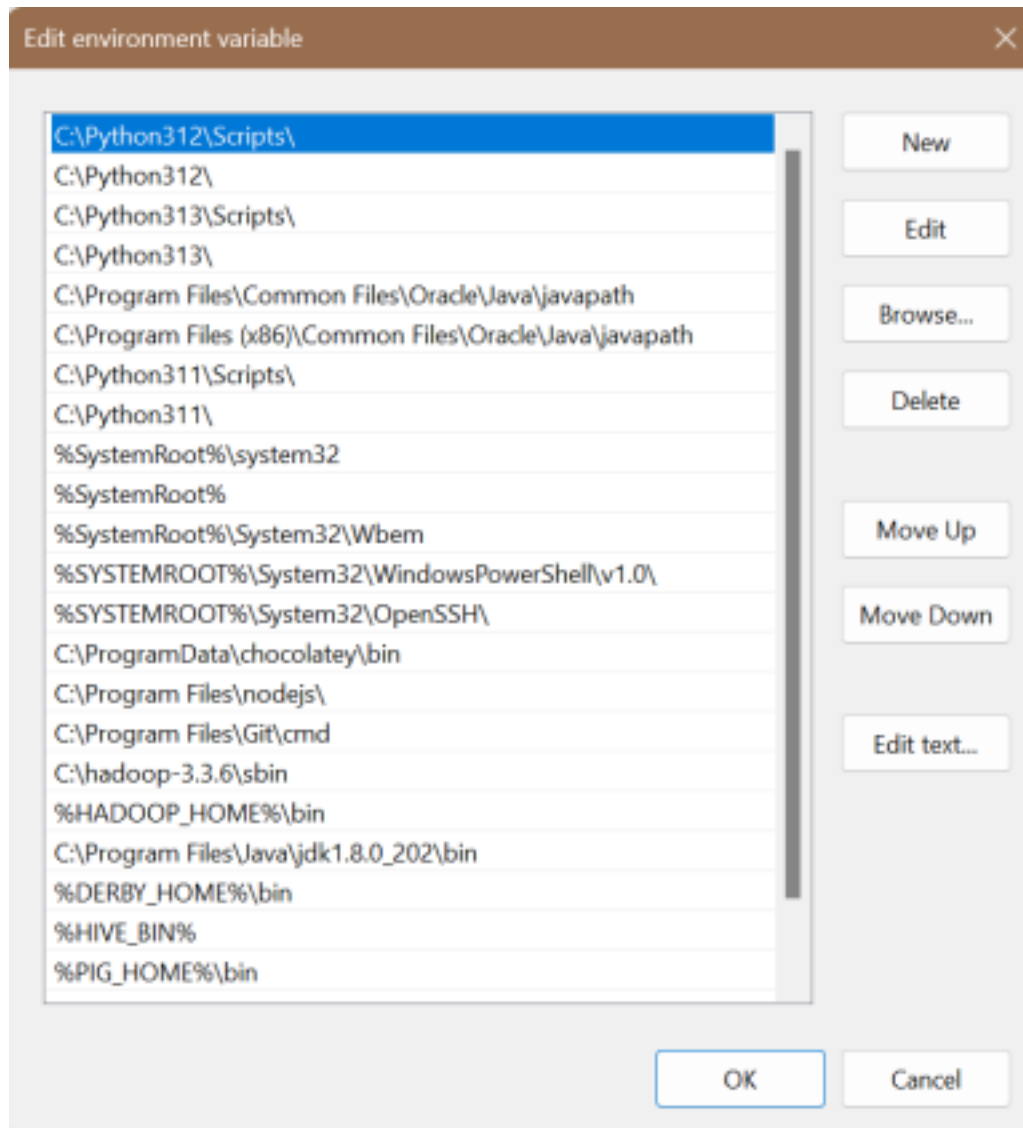
https://db.apache.org/derby/derby_downloads.html#For+Java+8+and+Higher

2. Downloading Apache Hive binaries

<https://downloads.apache.org/hive/hive-3.1.2/>

3. Setting environment variables

Variable	Value
HADOOP_HOME	C:\hadoop-3.3.6
HIVE_BIN	%HIVE_HOME%\bin
HIVE_HOME	C:\apache-hive-3.1.3-bin
HIVE_LIB	%HIVE_HOME%\lib
DERBY_HOME	C:\db-derby-10.14.2.0-bin



3.1. Copy Derby libraries

we should go to the Derby libraries directory (db-derby-10.14.2.0\lib) and copy all *.jar files. Then, we should paste them within the Hive libraries directory.

3.2. Configuring hive-site.xml and Hive's Bin folder

Refer following link to download the file. Also download the guava file. Put hive site.xml file to hive's conf location and replace hive's current guava file with this one in lib location. Also download the bin folder from link and replace the existing hive's bin folder.

<https://1drv.ms/f/s!ArSg3Xpur4Grmw0SDqW0g44T7HYU?e=wDsoBn>

4. Starting Hadoop Services

`start-all.cmd`

EX.NO.5 210701157

5. Derby Network Server: Run the following command in separate window to open Derby `StartNetworkServer -h 0.0.0.0`

6. Starting Apache Hive

Go to Apache Hive's bin location with cd command and run the following

command: `hive --service schematool -dbType derby -initSchema`

7. Open Hive shell by typing:

hive

Create a Database:

Start by creating a database. Open the Hive CLI and follow the steps

below: 1. Use the **CREATE DATABASE** statement to create a new

database: `CREATE DATABASE sample;`

2. Verify the database is present:

`SHOW DATABASES;`

3. Switch to the new database:

`USE sample;`

Create a Table in Hive

`CREATE TABLE students (name STRING, roll INT, dept STRING);`

Add Data

Run the **LOAD DATA LOCAL INPATH** command:

`LOAD DATA INPATH`

`'/C:/Users/mercy/OneDrive/Documents/DataAnalytics/Hive/students.csv' WRITE INTO
TABLE students;`

List Hive Tables and Data

To show all tables in a selected database, use the following statement:

`SHOW TABLES;`

To show table column names and data types, run:

`DESC students;`

EX.NO.5 210701157

To display table data, use a **SELECT** statement. For example, to select everything in a table, run:

`SELECT * FROM students;`

OUTPUT:

```
2024-08-26 20:02:33,376 INFO Hive:
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/C:/apache-hive-3.13.0/share/hadoop/common/lib/slf4j-reloadable-1.7.36.jar/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/C:/apache-hive-3.13.0/share/hadoop/common/lib/slf4j-reloadable-1.7.36.jar/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.hadoop.hive.log4j.HiveLoggerFactory]
2024-08-26 20:02:33,382 INFO conf.HiveConf: Found configuration file file:/C:/apache-hive-3.13.0-bin/conf/hive-site.xml
Hive Session ID = 6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:02:33,378 INFO SessionState: Hive Session ID = 6cbb0c57-f5d7-4a6e-9925-c62b4c956239

Logging initialized using configuration in jar:file:/C:/apache-hive-3.13.0-bin/lib/hive-common-3.13.0.jar!/hive-log4j.properties Async: true
2024-08-26 20:02:33,423 INFO SessionState:
Logging initialized using configuration in jar:file:/C:/apache-hive-3.13.0-bin/lib/hive-common-3.13.0.jar!/hive-log4j.properties Async: true
2024-08-26 20:02:34,869 INFO session.SessionState: Created HDFS directory: /tmp/hive/hwerc/6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:02:34,874 INFO session.SessionState: Created local directory: C:/Users/hwerc/AppData/Local/Temp/hwerc/6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:02:34,878 INFO session.SessionState: Created HDFS directory: /tmp/hive/hwerc/6cbb0c57-f5d7-4a6e-9925-c62b4c956239/_tmp_space.db
2024-08-26 20:02:34,888 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:02:34,888 INFO session.SessionState: Updating thread name to 6cbb0c57-f5d7-4a6e-9925-c62b4c956239 main
2024-08-26 20:02:34,790 INFO metastore.HiveMetaStore: 0: Opening new store with implementation class: org.apache.hadoop.hive.metastore.ObjectStore
2024-08-26 20:02:34,886 INFO metastore.ObjectStore: determineDefaultHiveMetaStore is set to unsupported value null. Setting it to value: ignored
2024-08-26 20:02:34,816 INFO metastore.ObjectStore: ObjectStore, initialize called
2024-08-26 20:02:34,812 INFO conf.HiveMetaStoreConf: Found configuration file file:/C:/apache-hive-3.13.0-bin/conf/hive-site.xml
2024-08-26 20:02:34,812 INFO conf.HiveMetaStoreConf: Unable to find config file HiveMetaStore-site.xml
2024-08-26 20:02:34,813 INFO conf.HiveMetaStoreConf: Found configuration file null
2024-08-26 20:02:34,813 INFO conf.HiveMetaStoreConf: Unable to find config file metastore-site.xml
2024-08-26 20:02:34,814 INFO conf.HiveMetaStoreConf: Found configuration file null
2024-08-26 20:02:34,942 INFO DataNucleus.Persistence: Property dataNucleus.cache.level1 unknown - will be ignored
2024-08-26 20:02:35,116 INFO Hikari.HikariDataSource: HikariPool-1 - Starting...
2024-08-26 20:02:35,614 INFO pool.PoolBase: HikariPool-1 - Driver does not support get/set network timeout for connections. (Feature not implemented: get/setNetworkTimeout.)
2024-08-26 20:02:35,713 INFO Hikari.HikariDataSource: HikariPool-1 - Start completed.
2024-08-26 20:02:35,707 INFO Hikari.HikariDataSource: HikariPool-2 - Starting...
```

```
hive> SHOW DATABASES;
2024-08-26 20:02:58,668 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:02:58,886 INFO ql.Driver: Compiling command(queryId=hwerc_20240826200159_8f9735af-41d8-4398-90d7-83f24e1a8032): SHOW DATABASES
2024-08-26 20:02:58,883 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:02:58,683 INFO ql.Driver: Semantic Analysis Completed (retry: false)
2024-08-26 20:02:58,730 INFO ql.Driver: Returning Hive schema: Schema(FieldSchemas:[FieldSchema{name:database_name, type:string, comment:from deserializer}], properties:null)
2024-08-26 20:02:58,829 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
2024-08-26 20:02:58,880 INFO ql.Driver: Completed compiling command(queryId=hwerc_20240826200159_8f9735af-41d8-4398-90d7-83f24e1a8032): Time taken: 0.77 seconds
2024-08-26 20:02:58,941 INFO exec.ReducibleDriver: Execution #1 of query
2024-08-26 20:02:58,846 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:02:58,886 INFO ql.Driver: Executing command(queryId=hwerc_20240826200159_8f9735af-41d8-4398-90d7-83f24e1a8032): SHOW DATABASES
2024-08-26 20:02:58,861 INFO ql.Driver: Starting task [Stage-0:000] in serial mode
2024-08-26 20:02:58,863 INFO metastore.HiveMetaStore: 0: get_databases: @hive#
2024-08-26 20:02:58,884 INFO HiveMetaStore.audit: ugi:hwerc ip=unknown-ip-addr cmd=get_databases: @hive#
2024-08-26 20:02:58,870 INFO exec.DDLTask: result: 2
2024-08-26 20:02:58,987 INFO ql.Driver: Completed executing command(queryId=hwerc_20240826200159_8f9735af-41d8-4398-90d7-83f24e1a8032): Time taken: 0.861 seconds
OK
2024-08-26 20:02:58,989 INFO ql.Driver: OK
2024-08-26 20:02:58,918 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:02:58,826 INFO conf.Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir
2024-08-26 20:02:58,961 INFO mapred.FileInputFormat: Total input files to process : 1
2024-08-26 20:02:58,982 INFO exec.ListSinkOperator: RECORD_OUT_INTERPRETERS:0, RECORD_OUT_OPERATOR:1,137_SINK_0:2,
default
sample
Time taken: 0.842 seconds, Fetched: 2 row(s)
2024-08-26 20:02:59,000 INFO CLIDriver: Time taken: 0.842 seconds, Fetched: 2 row(s)
2024-08-26 20:02:59,001 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:02:59,085 INFO session.SessionState: Resetting thread name to main
```

```
hive> USE sample
>
2024-08-26 20:05:59,786 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:05:59,795 INFO session.SessionState: Updating thread name to 6cbb0c57-f5d7-4a6e-9925-c62b4c956239 main
2024-08-26 20:05:59,799 INFO ql.Driver: Compiling command(queryId=hwerc_20240826200559_3d4f6251-d438-4a38-8429-0f648576237e): USE sample
2024-08-26 20:05:59,838 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:05:59,839 INFO metastore.HiveMetaStore: 0: get_database: @hive#sample
2024-08-26 20:05:59,821 INFO HiveMetaStore.audit: ugi:hwerc ip=unknown-ip-addr cmd=get_database: @hive#sample
2024-08-26 20:05:59,842 INFO ql.Driver: Semantic Analysis Completed (retry: false)
2024-08-26 20:05:59,842 INFO ql.Driver: Returning Hive schema: Schema(FieldSchemas:[null], properties:null)
2024-08-26 20:05:59,845 INFO ql.Driver: Completed compiling command(queryId=hwerc_20240826200559_3d4f6251-d438-4a38-8429-0f648576237e): Time taken: 0.845 seconds
2024-08-26 20:05:59,846 INFO exec.ReducibleDriver: Execution #1 of query
2024-08-26 20:05:59,888 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:05:59,851 INFO ql.Driver: Executing command(queryId=hwerc_20240826200559_3d4f6251-d438-4a38-8429-0f648576237e): USE sample
2024-08-26 20:05:59,852 INFO ql.Driver: Starting task [Stage-0:000] in serial mode
2024-08-26 20:05:59,883 INFO metastore.HiveMetaStore: 0: get_database: @hive#sample
2024-08-26 20:05:59,894 INFO HiveMetaStore.audit: ugi:hwerc ip=unknown-ip-addr cmd=get_database: @hive#sample
2024-08-26 20:05:59,888 INFO metastore.HiveMetaStore: 0: get_database: @hive#sample
2024-08-26 20:05:59,888 INFO HiveMetaStore.audit: ugi:hwerc ip=unknown-ip-addr cmd=get_database: @hive#sample
2024-08-26 20:05:59,883 INFO ql.Driver: Completed executing command(queryId=hwerc_20240826200559_3d4f6251-d438-4a38-8429-0f648576237e): Time taken: 0.813 seconds
OK
2024-08-26 20:05:59,864 INFO ql.Driver: OK
2024-08-26 20:05:59,864 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
Time taken: 0.866 seconds
2024-08-26 20:05:59,865 INFO CLIDriver: Time taken: 0.866 seconds
2024-08-26 20:05:59,865 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb0c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:05:59,886 INFO session.SessionState: Resetting thread name to main
```



```

hive> CREATE TABLE students (name STRING, roll INT, dept STRING);
2024-08-26 20:10:06,736 DFO conf.HiveConf: Using the default value passed in for log id: 6c8b1c57-f5d7-4a9e-9925-c62b4c956239
2024-08-26 20:10:06,736 DFO session.SessionState: Updating thread name to 6c8b1c57-f5d7-4a9e-9925-c62b4c956239 main
2024-08-26 20:10:06,736 DFO ql.Driver: Compiling command(queryId=mercy_20240826200908_74484479-9c70-487e-965b-ef577623a7de): CREATE TABLE students (name STRING, roll INT, dept STRING)
2024-08-26 20:10:06,780 DFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:10:06,784 DFO parse.CalcitePlanner: Starting Semantic Analysis
2024-08-26 20:10:06,880 DFO ugi.HiveMetaStoreClient: Created SQLNativeserviceController for session context : HiveAuthSessionContext {sessionIdString=6c8b1c57-f5d7-4a9e-9925-c62b4c956239, clientType=HIVECLIT}
2024-08-26 20:10:06,887 WARN session.SessionState: HETASTORE_FILTER_HOOK will be ignored, since hive.security.authorization.manager is set to instance of HiveAuthorizerFactory.
2024-08-26 20:10:06,888 DFO metastore.HiveMetaStoreClient: Metastore configuration metastore.filter.hook changed from org.apache.hadoop.hive.metastore.DefaultMetastoreFilterHook to org.apache.hadoop.hive.ql.security.authorization.plugin.AuthorizationMetastoreFilterHook
2024-08-26 20:10:06,889 DFO metastore.HiveMetaStore: 0: Cleaning up thread local RndStore...
2024-08-26 20:10:06,889 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=cleaning up thread local RndStore...
2024-08-26 20:10:06,890 DFO metastore.HiveMetaStore: 0: Done cleaning up thread local RndStore
2024-08-26 20:10:06,890 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=Done cleaning up thread local RndStore
2024-08-26 20:10:06,892 DFO metastore.HiveMetaStore: 0: Opening raw store with implementation class:org.apache.hadoop.hive.metastore.ObjectStore
2024-08-26 20:10:06,892 WARN metastore.ObjectStore: dataAccess.autoStartMechanismMode is set to unsupported value null . Setting it to value: ignored
2024-08-26 20:10:06,893 DFO metastore.ObjectStore: ObjectStore, Initialize called
2024-08-26 20:10:06,897 DFO metastore.MetastoreDirectSql: Using direct SQL, underlying DB is ORC9
2024-08-26 20:10:06,897 DFO metastore.ObjectStore: Initialized ObjectStore
2024-08-26 20:10:06,897 DFO metastore.RetryingMetastoreClient: RetryingMetastoreClient proxy class org.apache.hadoop.hive.ql.metastore.SessionHiveMetaStoreClient ugi=mercy (auth=SIMPLE) retried=1 delayed lifetime=0
2024-08-26 20:10:06,898 DFO parse.CalcitePlanner: Creating table sample.students position=1
2024-08-26 20:10:06,943 WARN metastore.ObjectStore: dataAccess.autoStartMechanismMode is set to unsupported value null . Setting it to value: ignored
2024-08-26 20:10:06,944 DFO metastore.ObjectStore: ObjectStore, Initialize called
2024-08-26 20:10:06,944 DFO metastore.MetastoreDirectSql: Using direct SQL, underlying DB is ORC9
2024-08-26 20:10:06,944 DFO metastore.ObjectStore: Initialized ObjectStore
2024-08-26 20:10:06,947 DFO metastore.RetryingMetastoreClient: RetryingMetastoreClient proxy class org.apache.hadoop.hive.ql.metastore.SessionHiveMetaStoreClient ugi=mercy (auth=SIMPLE) retried=1 delayed lifetime=0
2024-08-26 20:10:06,952 DFO metastore.HiveMetaStore: 0: get_database: @hive/sample
2024-08-26 20:10:06,956 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_database: @hive/sample
2024-08-26 20:10:06,977 DFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-08-26 20:10:06,978 DFO ql.Driver: Returning Hive schema: Schema(FieldSchemas:null, properties:null)
2024-08-26 20:10:06,978 DFO ql.Driver: Completed compiling command(queryId=mercy_20240826200908_74484479-9c70-487e-965b-ef577623a7de); Time taken: 0.35 seconds
2024-08-26 20:10:06,982 DFO rexec.RexecDriver: Execution #1 of query
2024-08-26 20:10:06,983 DFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:10:06,987 DFO ql.Driver: Executing command(queryId=mercy_20240826200908_74484479-9c70-487e-965b-ef577623a7de): CREATE TABLE students (name STRING, roll INT, dept STRING)
2024-08-26 20:10:06,995 DFO ql.Driver: Starting task [Stage-0:ORC] in serial mode
2024-08-26 20:10:06,996 DFO metastore.HiveMetaStoreClient: Metastore configuration metastore.filter.hook changed from org.apache.hadoop.hive.ql.security.authorization.plugin.AuthorizationMetastoreFilterHook to org.apache.hadoop.hive.metastore.DefaultMetastoreFilterHook
2024-08-26 20:10:06,996 DFO metastore.HiveMetaStore: 0: Cleaning up thread local RndStore...

```

EX.NO.5

210701157

```

hive> LOAD DATA LOCAL INPATH '/C:/Users/mercy/OneDrive/Documents/DataAnalytics/Hive/students.csv' OVERWRITE INTO TABLE students;
2024-08-26 20:10:06,846 DFO conf.HiveConf: Using the default value passed in for log id: 6c8b1c57-f5d7-4a9e-9925-c62b4c956239
2024-08-26 20:10:06,880 DFO session.SessionState: Updating thread name to 6c8b1c57-f5d7-4a9e-9925-c62b4c956239 main
2024-08-26 20:10:06,889 DFO ql.Driver: Compiling command(queryId=mercy_20240826200908_7518333e-e17b-45ab-9b6d-7ec2e9646e6a): LOAD DATA LOCAL INPATH '/C:/Users/mercy/OneDrive/Documents/DataAnalytics/Hive/students.csv' OVERWRITE INTO TABLE students
2024-08-26 20:10:06,887 DFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:10:06,887 DFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-26 20:10:06,890 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-26 20:10:06,875 DFO io.HiveFileFormatUtil: Failed to check mode for C:/Users/mercy/OneDrive/Documents/DataAnalytics/Hive/students.csv: View 'unix' not available (class java.lang.UnsupportedOperationException)
2024-08-26 20:10:06,898 DFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-08-26 20:10:06,899 DFO ql.Driver: Returning Hive schema: Schema(FieldSchemas:null, properties:null)
2024-08-26 20:10:06,899 DFO ql.Driver: Completed compiling command(queryId=mercy_20240826200908_7518333e-e17b-45ab-9b6d-7ec2e9646e6a); Time taken: 0.293 seconds
2024-08-26 20:10:06,900 DFO rexec.RexecDriver: Execution #1 of query
2024-08-26 20:10:06,941 DFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:10:06,941 DFO ql.Driver: Executing command(queryId=mercy_20240826200908_7518333e-e17b-45ab-9b6d-7ec2e9646e6a): LOAD DATA LOCAL INPATH '/C:/Users/mercy/OneDrive/Documents/DataAnalytics/Hive/students.csv' OVERWRITE INTO TABLE students
2024-08-26 20:10:06,942 DFO ql.Driver: Starting task [Stage-0:HIVE] in serial mode
2024-08-26 20:10:06,942 DFO metastore.HiveMetaStore: 0: Cleaning up thread local RndStore...
2024-08-26 20:10:06,942 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=cleaning up thread local RndStore...
2024-08-26 20:10:06,943 DFO metastore.HiveMetaStore: 0: Done cleaning up thread local RndStore
2024-08-26 20:10:06,943 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=Done cleaning up thread local RndStore
Loading data to table sample.students
2024-08-26 20:10:06,994 DFO exec.Task: Loading data to table sample.students from file:/C:/Users/mercy/OneDrive/Documents/DataAnalytics/Hive/students.csv
2024-08-26 20:10:06,945 DFO metastore.HiveMetaStore: 0: Opening raw store with implementation class:org.apache.hadoop.hive.metastore.ObjectStore
2024-08-26 20:10:06,946 WARN metastore.ObjectStore: dataAccess.autoStartMechanismMode is set to unsupported value null . Setting it to value: ignored
2024-08-26 20:10:06,946 DFO metastore.ObjectStore: ObjectStore, Initialize called
2024-08-26 20:10:06,950 DFO metastore.MetastoreDirectSql: Using direct SQL, underlying DB is ORC9
2024-08-26 20:10:06,950 DFO metastore.ObjectStore: Initialized ObjectStore
2024-08-26 20:10:06,951 DFO metastore.RetryingMetastoreClient: RetryingMetastoreClient proxy class org.apache.hadoop.hive.ql.metastore.SessionHiveMetaStoreClient ugi=mercy (auth=SIMPLE) retried=1 delayed lifetime=0
2024-08-26 20:10:06,951 DFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-26 20:10:06,952 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-26 20:10:06,986 DFO io.HiveFileFormatUtil: Failed to check mode for C:/Users/mercy/OneDrive/Documents/DataAnalytics/Hive/students.csv: View 'unix' not available (class java.lang.UnsupportedOperationException)
2024-08-26 20:10:06,989 DFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-26 20:10:06,990 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-26 20:10:06,123 DFO metastore.HiveMetaStore: 0: get_database: @hive/sample
2024-08-26 20:10:06,124 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_database: @hive/sample
2024-08-26 20:10:06,133 DFO common.FileUtil: Creating directory if it doesn't exist: hdfs://localhost:9000/user/hive/warehouse/sample.db/students
2024-08-26 20:10:06,946 DFO metastore.HiveMetaStore: 0: alter_table : hive.sample.students newtbl=students
2024-08-26 20:10:06,967 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=alter_table: hive.sample.students newtbl=students
2024-08-26 20:10:18,885 DFO ql.Driver: Starting task [Stage-1:STATS] in serial mode
2024-08-26 20:10:18,886 DFO metastore.HiveMetaStore: 0: Cleaning up thread local RndStore...
2024-08-26 20:10:18,886 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=cleaning up thread local RndStore...
hive> SHOW TABLES;
2024-08-26 20:40:06,827 DFO conf.HiveConf: Using the default value passed in for log id: 6c8b1c57-f5d7-4a9e-9925-c62b4c956239
2024-08-26 20:40:06,828 DFO session.SessionState: Updating thread name to 6c8b1c57-f5d7-4a9e-9925-c62b4c956239 main
2024-08-26 20:40:06,836 DFO ql.Driver: Compiling command(queryId=mercy_20240826204006_8d158d3-c6f2-47b4-84f6-7381c6428533): SHOW TABLES
2024-08-26 20:40:06,865 DFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:40:06,868 DFO metastore.HiveMetaStore: 0: get_database: @hive/sample
2024-08-26 20:40:06,868 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_database: @hive/sample
2024-08-26 20:40:06,873 DFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-08-26 20:40:06,874 DFO ql.Driver: Returning Hive schema: Schema(FieldSchemas:[FieldSchema{name=tbl_name, type=string, comment=from deserializer}], properties:null)
2024-08-26 20:40:06,876 DFO exec.LimitedOperator: Initializing operator LST_53W6[0]
2024-08-26 20:40:06,876 DFO ql.Driver: Completed compiling command(queryId=mercy_20240826204006_8d158d3-c6f2-47b4-84f6-7381c6428533); Time taken: 0.848 seconds
2024-08-26 20:40:06,876 DFO rexec.RexecDriver: Execution #0 of query
2024-08-26 20:40:06,877 DFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:40:06,877 DFO ql.Driver: Executing command(queryId=mercy_20240826204006_8d158d3-c6f2-47b4-84f6-7381c6428533): SHOW TABLES
2024-08-26 20:40:06,878 DFO ql.Driver: Starting task [Stage-0:TEXT] in serial mode
2024-08-26 20:40:06,880 DFO metastore.HiveMetaStore: 0: get_database: @hive/sample
2024-08-26 20:40:06,881 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_database: @hive/sample
2024-08-26 20:40:06,885 DFO metastore.HiveMetaStore: 0: get_tables: db=@hive/sample path=.
2024-08-26 20:40:06,886 DFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_tables: db=@hive/sample path=.
2024-08-26 20:40:06,916 DFO ql.Driver: Completed executing command(queryId=mercy_20240826204006_8d158d3-c6f2-47b4-84f6-7381c6428533); Time taken: 0.841 seconds
hive>
2024-08-26 20:40:06,916 DFO ql.Driver: OK
2024-08-26 20:40:06,926 DFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:40:06,924 DFO mapred.FileInputFormat: Total input files to process : 1
2024-08-26 20:40:06,954 DFO exec.LimitedOperator: RECORDS_OUT_INTERMEDIATE=0, RECORDS_OUT_OPERATOR_LIST_SIZE=0,1,
students
Time taken: 0.691 seconds, Fetched: 1 row(s)
2024-08-26 20:40:06,963 DFO CLIDriver: Time taken: 0.691 seconds, Fetched: 1 row(s)
2024-08-26 20:40:06,963 DFO conf.HiveConf: Using the default value passed in for log id: 6c8b1c57-f5d7-4a9e-9925-c62b4c956239
2024-08-26 20:40:06,964 DFO session.SessionState: Resetting thread name to main
hive>

```

```

hive> DESC students;
2024-08-26 20:48:36,768 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb1c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:48:36,769 INFO session.SessionState: Updating thread name to 6cbb1c57-f5d7-4a6e-9925-c62b4c956239 main
2024-08-26 20:48:36,770 INFO ql.Driver: Compiling command(queryId=mercy_20240826204836_742eacea-f18c-4465-81ba-0c86b0bbaeae2): DESC students
2024-08-26 20:48:36,783 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:48:36,784 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-26 20:48:36,786 INFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-26 20:48:36,831 INFO parse.SOLSemanticAnalyzer: analyzeDescribeTable done
2024-08-26 20:48:36,831 INFO ql.Driver: Semantic Analysis Completed (retryId = false)
2024-08-26 20:48:36,832 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:col_name, type:string, comment:from deserializer), FieldSchema(n_type, type:string, comment:from deserializer), FieldSchema(name:comment, type:string, comment:from deserializer)], properties:null)
2024-08-26 20:48:36,834 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
2024-08-26 20:48:36,834 INFO ql.Driver: Completed compiling command(queryId=mercy_20240826204836_742eacea-f18c-4465-81ba-0c86b0bbaeae2); Time taken: 0.064 seconds
2024-08-26 20:48:36,834 INFO exec.RefExecDriver: Execution #1 of query
2024-08-26 20:48:36,834 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:48:36,834 INFO ql.Driver: Executing command(queryId=mercy_20240826204836_742eacea-f18c-4465-81ba-0c86b0bbaeae2): DESC students
2024-08-26 20:48:36,834 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-08-26 20:48:36,838 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-26 20:48:36,838 INFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-26 20:48:36,873 INFO ql.Driver: Completed executing command(queryId=mercy_20240826204836_742eacea-f18c-4465-81ba-0c86b0bbaeae2); Time taken: 0.037 seconds
OK
2024-08-26 20:48:36,873 INFO ql.Driver: OK
2024-08-26 20:48:36,873 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:48:36,875 INFO mapped.FileInputFormat: Total input files to process : 1
2024-08-26 20:48:36,896 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:3,
name
string
roll
int
dept
string
Time taken: 0.289 seconds, Fetched: 3 row(s)
2024-08-26 20:48:36,915 INFO CLIDriver: Time taken: 0.189 seconds, Fetched: 3 row(s)
2024-08-26 20:48:36,917 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb1c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:48:36,923 INFO session.SessionState: Resetting thread name to main
hive>

```

```

hive> SELECT * FROM students;
2024-08-26 20:41:12,643 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb1c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:41:12,644 INFO session.SessionState: Updating thread name to 6cbb1c57-f5d7-4a6e-9925-c62b4c956239 main
2024-08-26 20:41:12,645 INFO ql.Driver: Compiling command(queryId=mercy_20240826204112_4f6f657c-0e35-400c-9a40-f9a19709922d): SELECT * F
2024-08-26 20:41:12,663 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:41:12,663 INFO parse.CalcitePlanner: Starting Semantic Analysis
2024-08-26 20:41:12,691 INFO parse.CalcitePlanner: Completed phase 1 of Semantic Analysis
2024-08-26 20:41:12,691 INFO parse.CalcitePlanner: Get metadata for source tables
2024-08-26 20:41:12,691 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.sample.students
2024-08-26 20:41:12,693 INFO HiveMetaStore.audit: ugi=mercy ip=unknown-ip-addr cmd=get_table : tbl=hive.sample.students
2024-08-26 20:41:12,751 INFO parse.CalcitePlanner: Get metadata for subqueries
2024-08-26 20:41:12,751 INFO parse.CalcitePlanner: Get metadata for destination tables
2024-08-26 20:41:12,791 INFO ql.Context: New scratch dir is hdfs://localhost:9000/tmp/hive/mercy/6cbb1c57-f5d7-4a6e-9925-c62b4c956239/hj
9444333444-1
2024-08-26 20:41:13,783 INFO parse.CalcitePlanner: Completed getting MetaData for Semantic Analysis
OK
2024-08-26 20:41:14,144 INFO ql.Driver: OK
2024-08-26 20:41:14,145 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-08-26 20:41:14,160 INFO mapped.FileInputFormat: Total input files to process : 1
2024-08-26 20:41:14,288 INFO lazy.LazyStruct: Missing fields! Expected 3 fields but only got 1! Ignoring similar problems.
2024-08-26 20:41:14,291 INFO exec.TableScanOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_T5_0:10,
2024-08-26 20:41:14,291 INFO exec.SelectOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_SEL_1:10,
2024-08-26 20:41:14,291 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_3:10,
name,roll,dept NULL NULL
Abhishek,110,CSE NULL NULL
Pragya,126,CSE NULL NULL
Dev, 231,CSEB NULL NULL
Gonakshi,241,CSEB NULL NULL
Sameer,343,IT NULL NULL
Naina,356,IT NULL NULL
Ranveer,511,ECE NULL NULL
Ishaani,521,ECE NULL NULL
Mercy,157,CSE NULL NULL
Time taken: 1.5 seconds, Fetched: 10 row(s)
2024-08-26 20:41:14,321 INFO CLIDriver: Time taken: 1.5 seconds, Fetched: 10 row(s)
2024-08-26 20:41:14,321 INFO conf.HiveConf: Using the default value passed in for log id: 6cbb1c57-f5d7-4a6e-9925-c62b4c956239
2024-08-26 20:41:14,322 INFO session.SessionState: Resetting thread name to main
hive>

```

RESULT:

Thus a table has been created in Hive and writing queries to access the data in the table is completed successfully.